

版权声明

这些幻灯片根据 Creative Commons License 分发。

DeepLearning.AI 将这些幻灯片用于教育目的。您不得出于商业目的使用或分发这些幻灯片。您可以复制这些幻灯片并将其用于教育目的，只要您引用 DeepLearning.AI 作为幻灯片的来源即可。

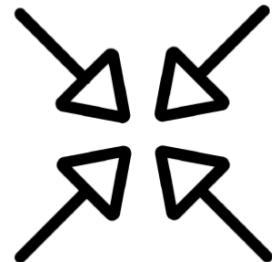
有关许可证的其余详细信息，请参阅
<https://creativecommons.org/licenses/by-sa/2.0/legalcode>



模式折叠

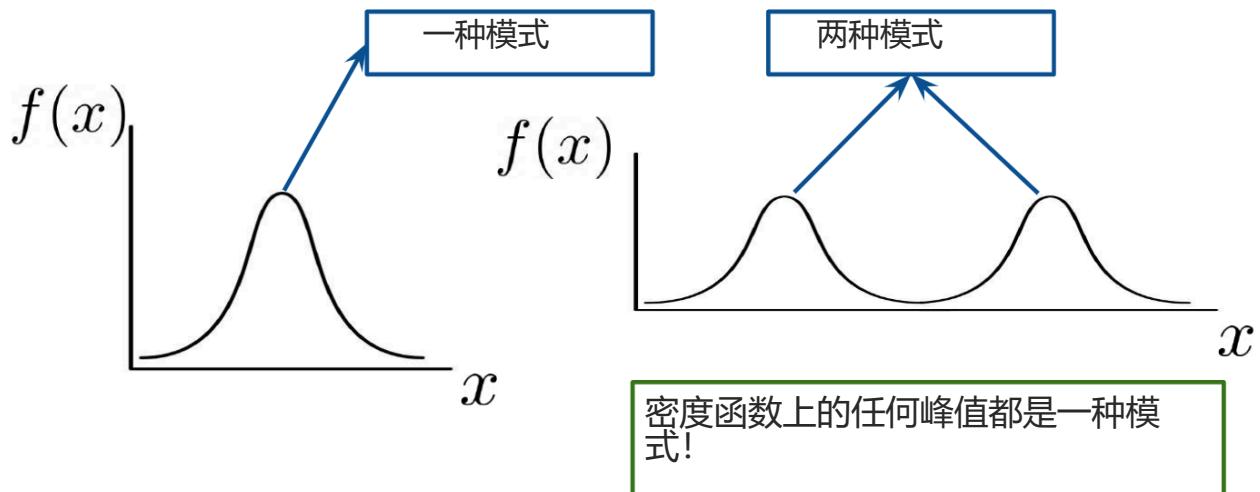
大纲

- 发行版中的模式
- GAN 中的模式崩溃
- 训练期间的直觉



© deeplearning.ai

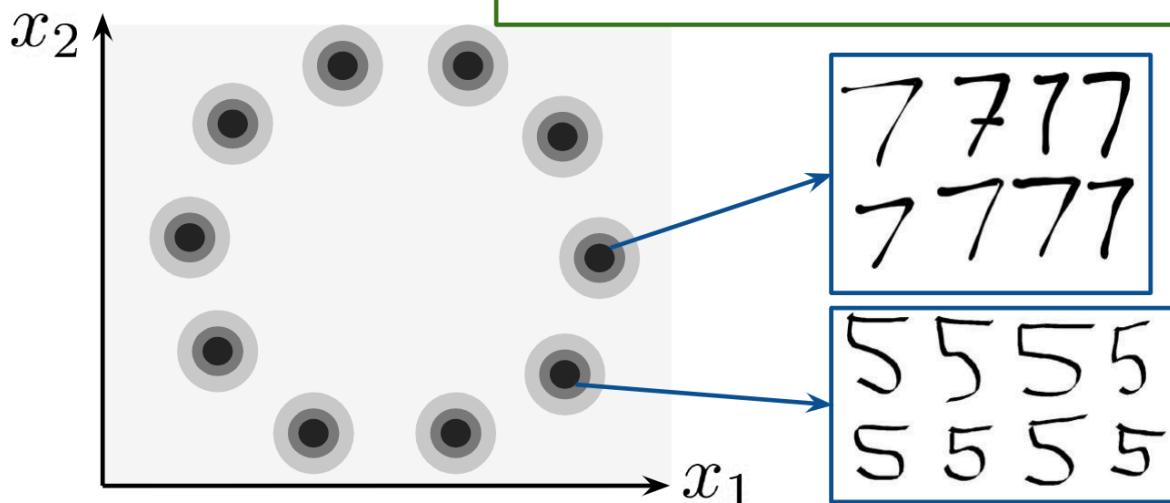
模式折叠



© deeplearning.ai

模式折叠

10 种不同的模式，每位数字 1 个



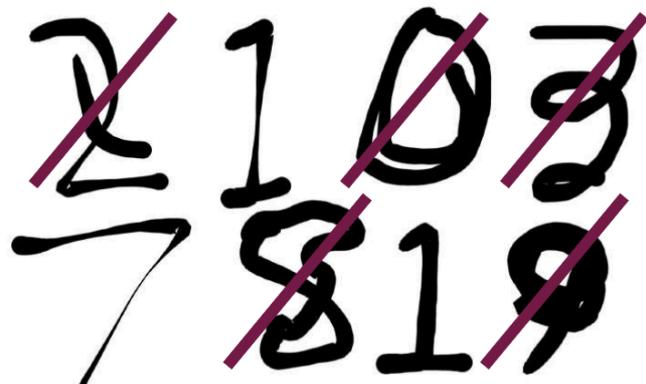
© deeplearning.ai

模式折叠



© deeplearning.ai

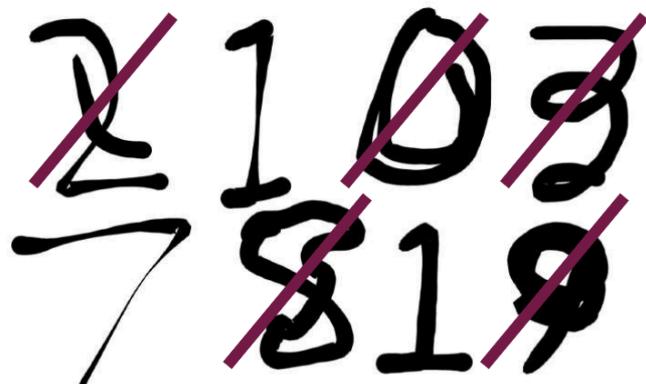
模式折叠



假货

© deeplearning.ai

模式折叠

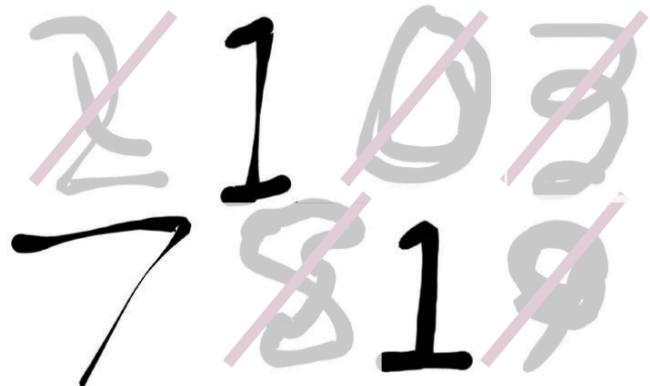


© deeplearning.ai

模式折叠



发电机



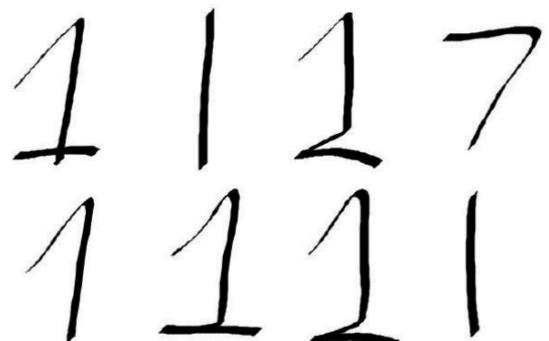
欺骗鉴别器的
假货

© deeplearning.ai

模式折叠



发电机



© deeplearning.ai

模式折叠



1 1 1 1
1 1 1 1

假货

© deeplearning.ai

模式折叠



1 1 1 1
1 1 1 1

© deeplearning.ai

模式折叠



发电机

1 1 1
1 1 1 1

欺骗鉴别器的
假货

© deeplearning.ai

模式折叠



发电机

1 1 1 1
1 1 1 1

© deeplearning.ai

总结

- 众数是特征分布中的峰值
- 真实数据集的典型情况
- 当生成器卡在一种模式下时，会发生模式崩溃



© deeplearning.ai

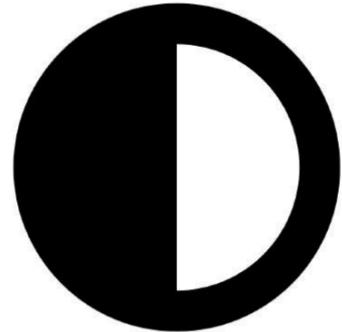


deeplearning.ai

问题
BCE 损失

大纲

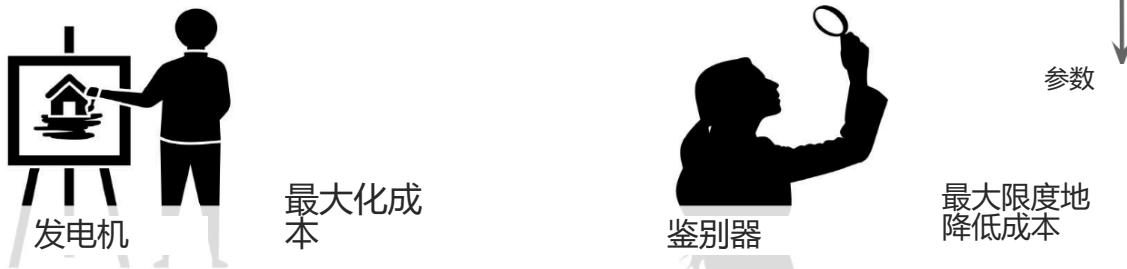
- GAN 中的 BCE 损失和最终目标
- BCE 损失问题



© deeplearning.ai

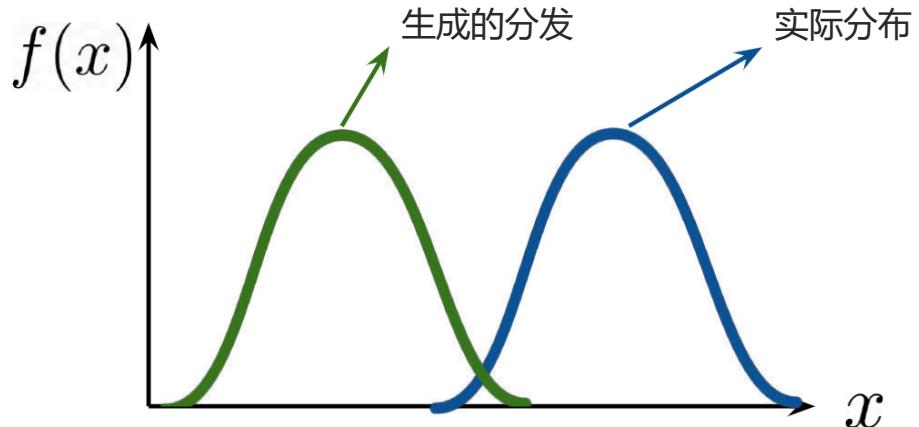
GAN 中的 BCE 损失

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$



© deeplearning.ai

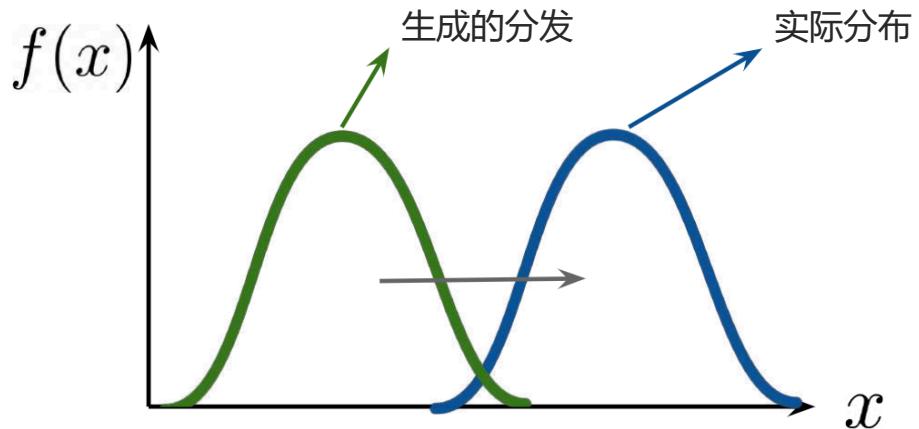
GAN 中的目标



© deeplearning.ai

GAN 中的目标

使生成的分布和实际分布看起来相似



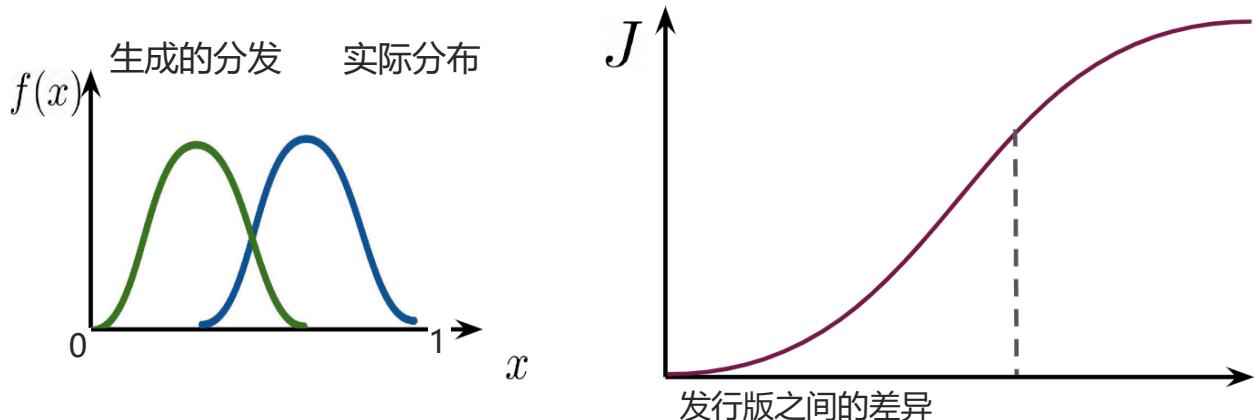
© deeplearning.ai

GAN 中的 BCE 损失批评更直接



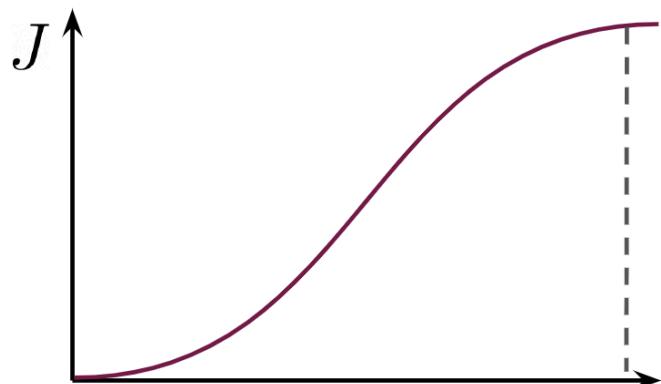
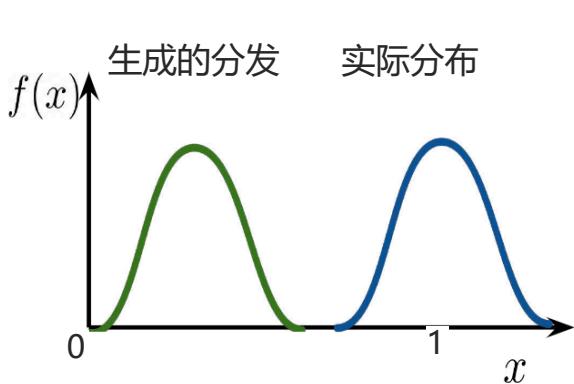
© deeplearning.ai

BCE Loss 的问题



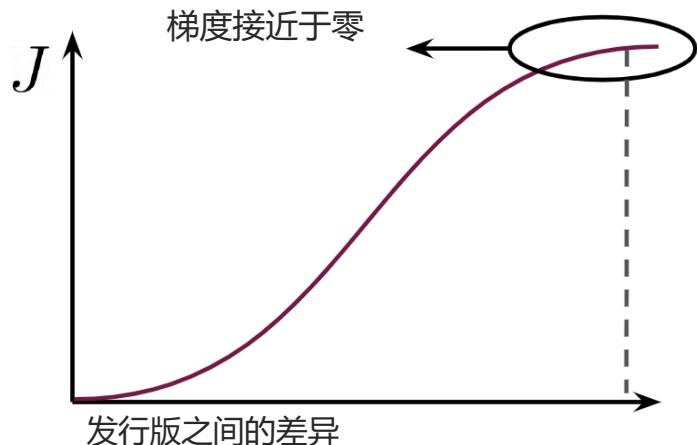
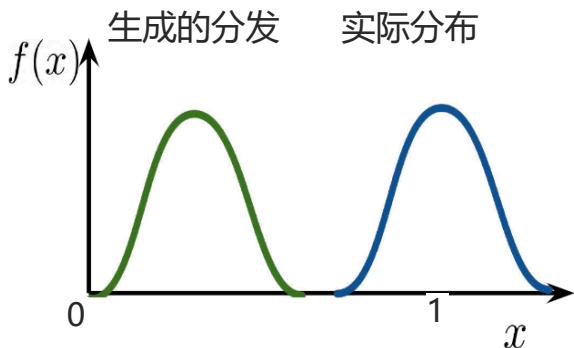
© deeplearning.ai

BCE Loss 的问题



© deeplearning.ai

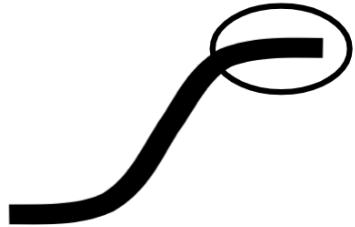
BCE Loss 的问题



© deeplearning.ai

总结

- GAN 尝试使真实分布和生成的分布看起来相似
- 当判别器改进太多时，由 BCE Loss 近似的函数将包含平坦区域
- 成本函数上的平坦区域 = 梯度消失



© deeplearning.ai



Earth Mover的
距离

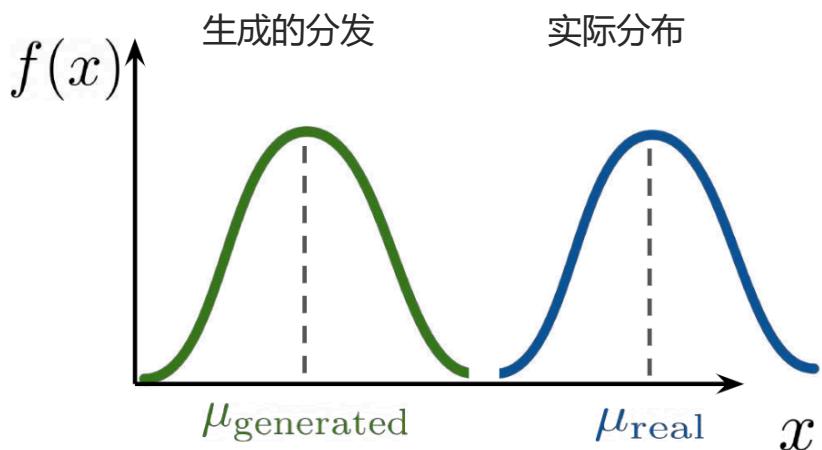
大纲

- 推土机距离 (EMD)
- 为什么它解决了 BCE Loss 的梯度消失问题



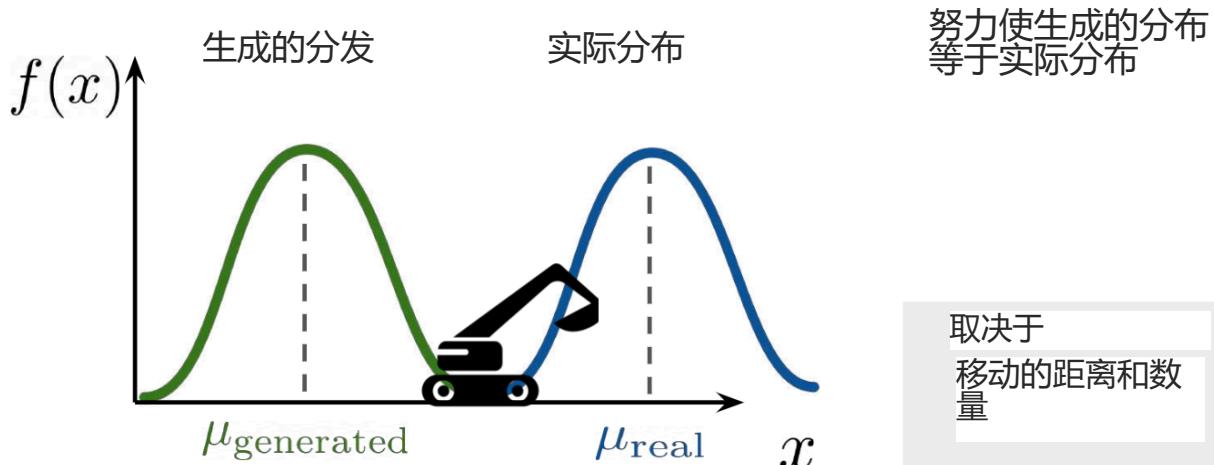
© deeplearning.ai

推土机的距离



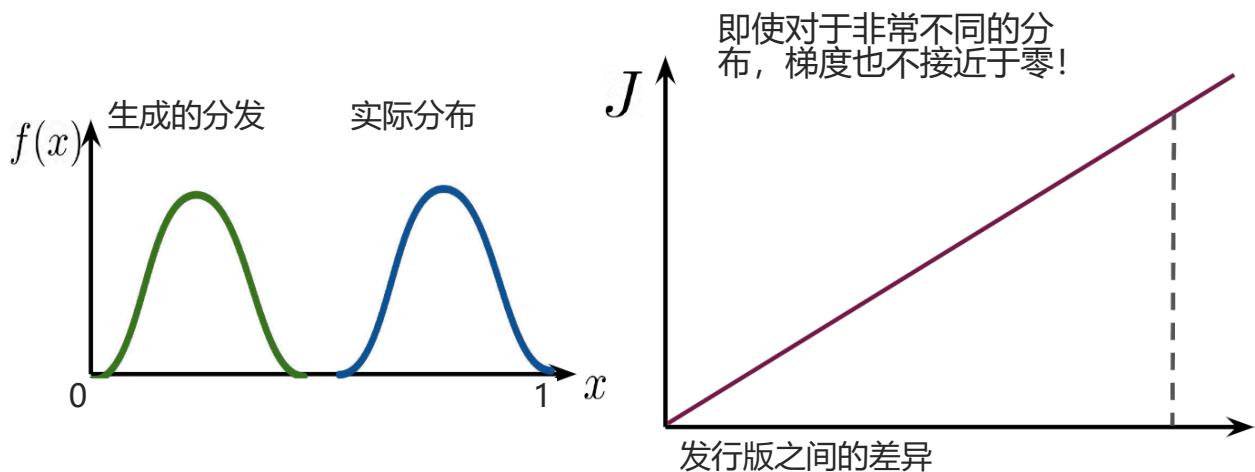
© deeplearning.ai

推土机的距离



© deeplearning.ai

推土机的距离



© deeplearning.ai

总结

- 推土机距离 (EMD) 是数量和距离的函数
- 当分布差异很大时，没有平坦的区域
- 近似 EMD 解决了与 BCE 相关的问题



© deeplearning.ai



deeplearning.ai

水垢
损失

大纲

- 简化的 BCE 损失
- W-Loss 及其与 BCE Loss 的比较



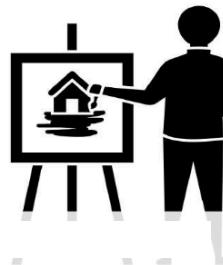
© deeplearning.ai

简化的 BCE 损失

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$



最大限度地
发电机



最大化成
本

© deeplearning.ai

简化的 BCE 损失

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

$$\min_d \max_g -[\mathbb{E}(\log(d(x))) + \mathbb{E}()]$$



最大限度地降低成本



最大化成本

© deeplearning.ai

简化的 BCE 损失

$$J(\theta) = -\frac{1}{m} \sum_{i=1}^m [y^{(i)} \log h(x^{(i)}, \theta) + (1 - y^{(i)}) \log(1 - h(x^{(i)}, \theta))]$$

$$\min_d \max_g -[\mathbb{E}(\log(d(x))) + \mathbb{E}(1 - \log(d(g(z))))]$$



最大限度地降低成本

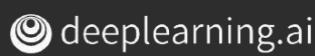


最大化成本

© deeplearning.ai

W 损失

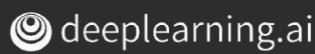
W-Loss 近似于地球移动器的距离



W 损失

W-Loss 近似于地球移动器的距离

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$$



W 损失

W-Loss 近似于地球移动器的距离

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$$



 deplearning.ai

W 损失

W-Loss 近似于地球移动器的距离

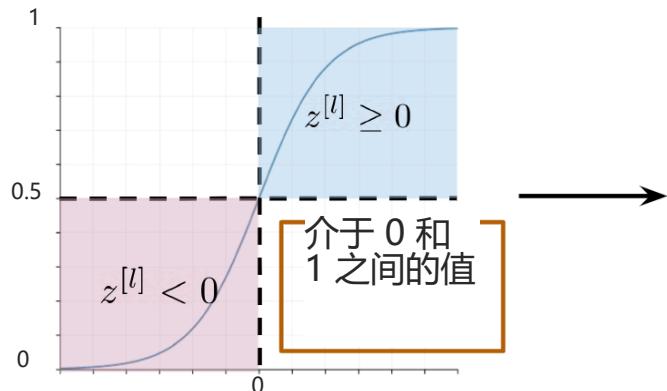
$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$$



 deplearning.ai

甄别器输出

鉴别器输出



鉴别器输出

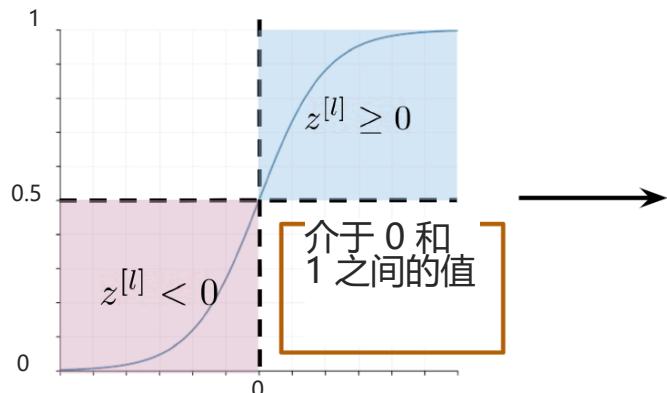
任何 Real
价值

deeplearning.ai

甄别器输出

鉴别器输出

鉴别器输出
评论家



任何真实的
价值

deeplearning.ai

W 损失与 BCE 损失

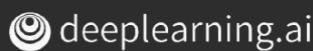
BCE 损失	W 损失
鉴别器输出介于 0 和 1 之间	Critic 输出任意数字
$-\left[\mathbb{E}(\log(d(x))) + \mathbb{E}(1 - \log(d(g(z))))\right]$	$\mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$

W-Loss 有助于解决模式崩溃和梯度消失问题



总结

- W-Loss 看起来与 BCE Loss 非常相似
- W-Loss 可防止模式崩溃和梯度消失问题





Wasserstein Critic 的条件

大纲

- 评论家神经网络上的连续性条件
- 为什么这种情况很重要



W 损失的条件

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$$

© deeplearning.ai

W 损失的条件

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$$

© deeplearning.ai

W 损失的条件

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))$$

© deeplearning.ai

W 损失的条件

$$\min_g \max_c \mathbb{E}(\boxed{c}(x)) - \mathbb{E}(\boxed{c}(g(z)))$$

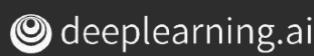
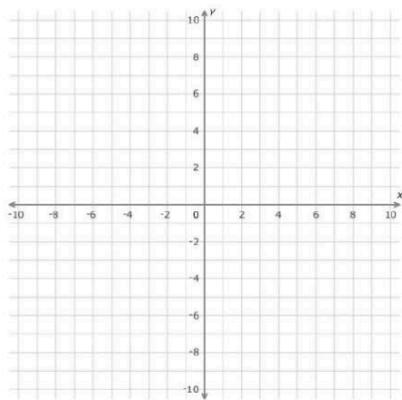
需要为 1-Lipschitz 连续

© deeplearning.ai

W 损失的条件

Critic 需要为 1 L 连续

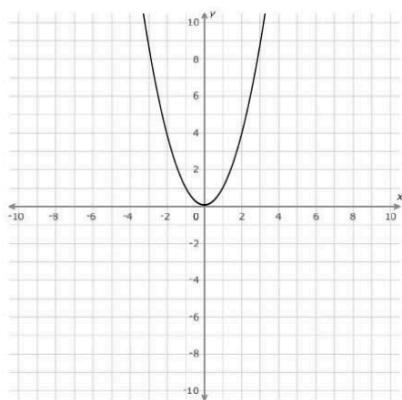
对于每个点，梯度的范数应最多为 1



W 损失的条件

Critic 需要为 1 L 连续

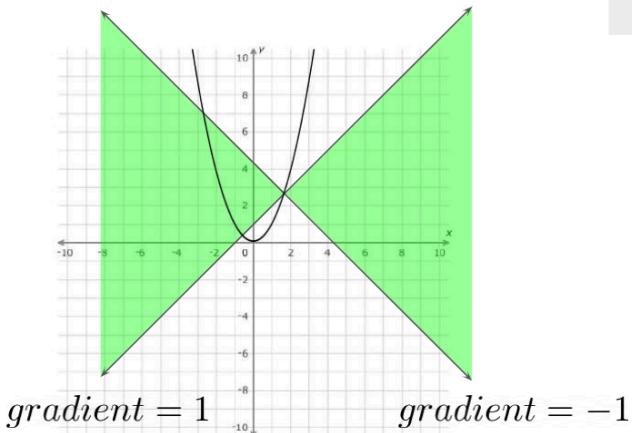
对于每个点，梯度的范数应最多为 1



W 损失的条件

Critic 需要为 1 L 连续

对于每个点，梯度的范数应最多为 1

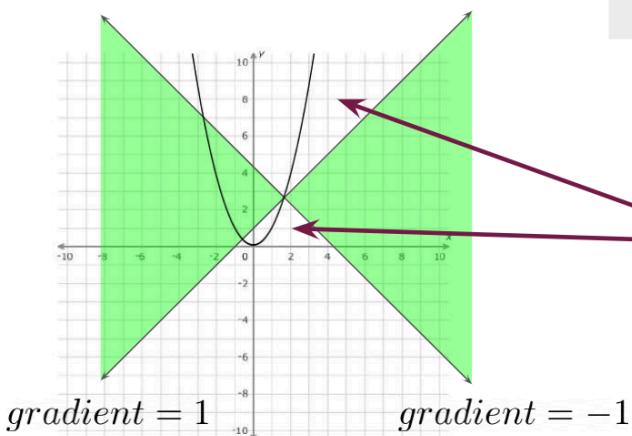


deeplearning.ai

W 损失的条件

Critic 需要为 1 L 连续

梯度的范数应为
每个点最多 1 个



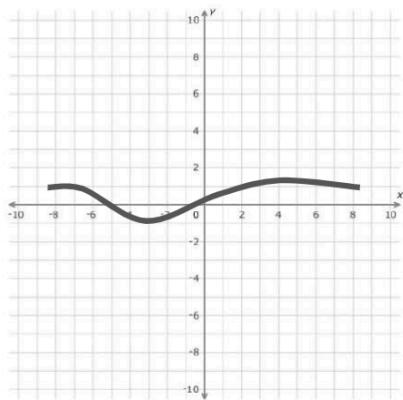
不是 1 L 连续

deeplearning.ai

W 损失的条件

Critic 需要为 1 L 连续

对于每个点，梯度的范数应最多为 1

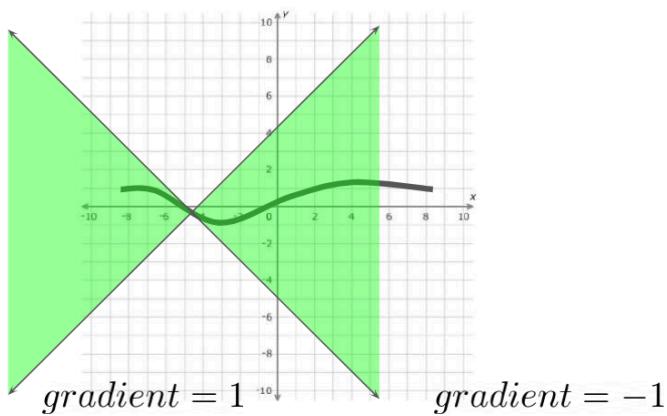


 deplearning.ai

W 损失的条件

Critic 需要为 1 L 连续

对于每个点，梯度的范数应最多为 1

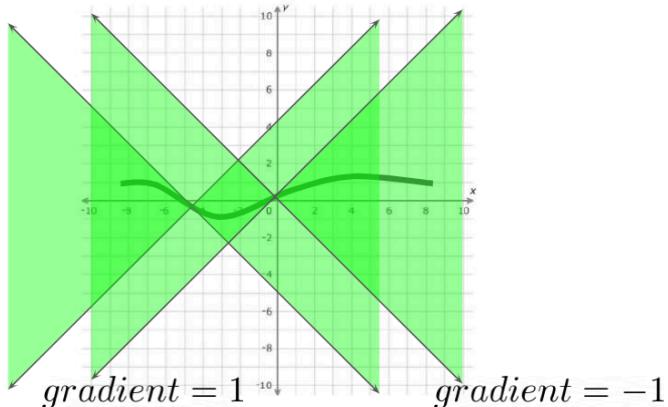


 deplearning.ai

W 损失的条件

Critic 需要为 1 L 连续

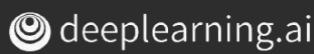
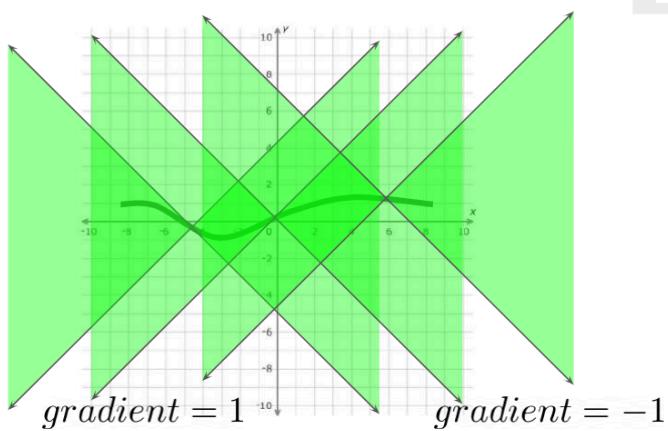
对于每个点，梯度的范数应最多为 1



W 损失的条件

Critic 需要为 1 L 连续

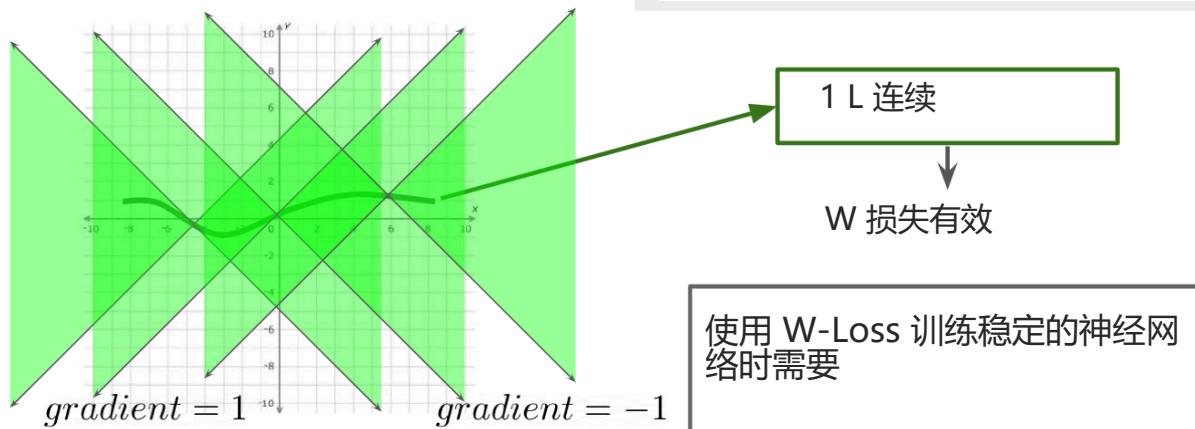
对于每个点，梯度的范数应最多为 1



W 损失的条件

Critic 需要为 1 L 连续

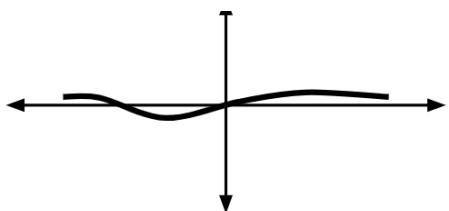
对于每个点，梯度的范数应最多为 1



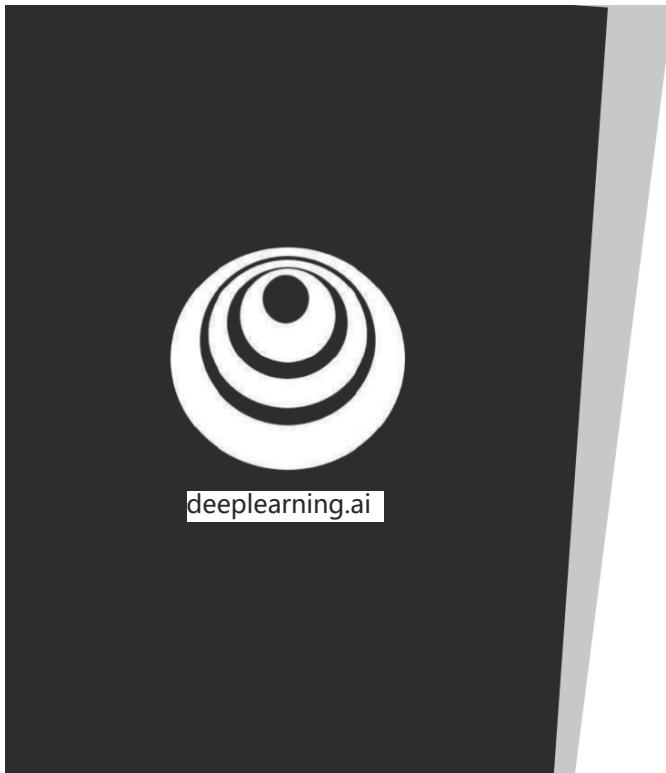
deeplearning.ai

总结

- 使用 W-Loss 时，Critic 的神经网络需要为 1-L Continuous
- 此条件可确保 W-Loss 有效地近似地接近地球移动器的距离



deeplearning.ai



1-Lipschitz 连续性执行

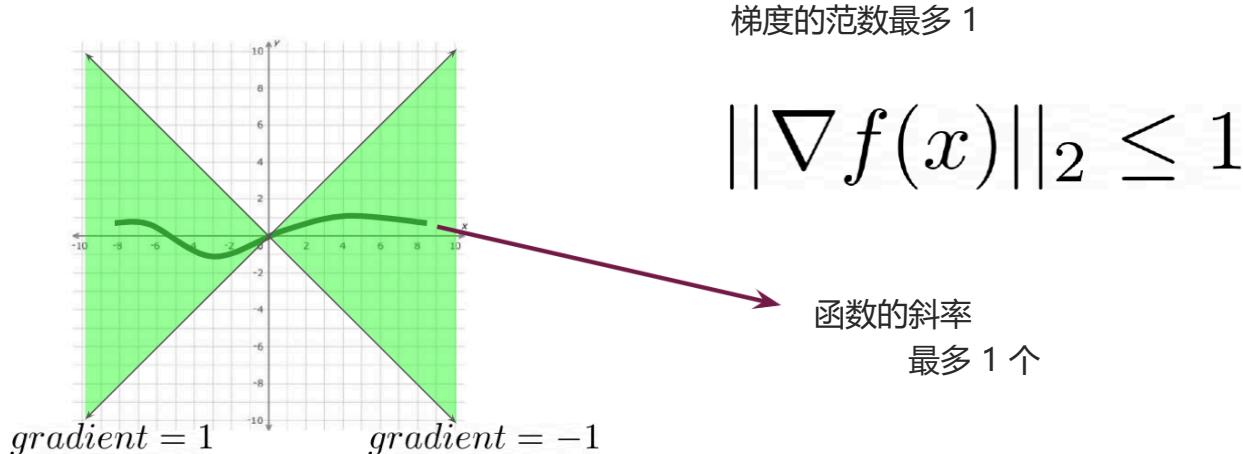
大纲

- 权重裁剪和渐变惩罚
- 梯度惩罚的优点



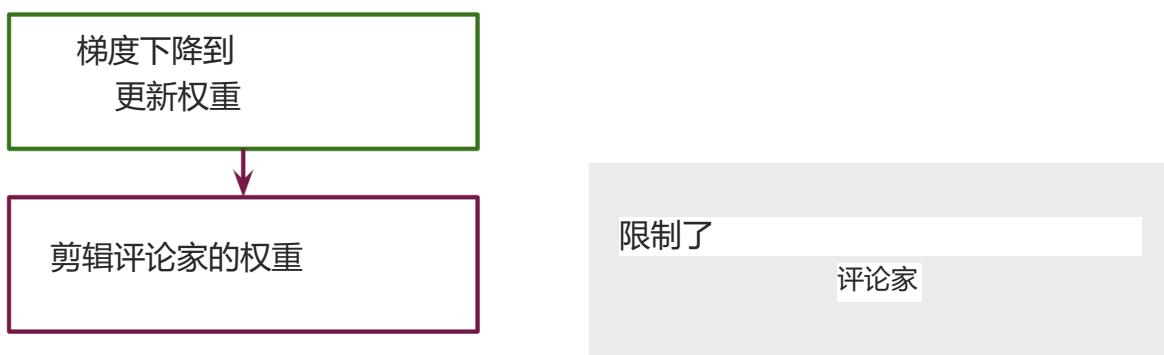
1-L 强制执行

Critic 需要为 1 L 连续



1 L 强制：重量夹

权重剪切将评论家的权重强制到固定的间隔



1-L 执行: 梯度惩罚

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) + \lambda \text{reg}$$

↓
 正则化
 评论家的渐变

© deeplearning.ai

1-L 执行: 梯度惩罚

真正



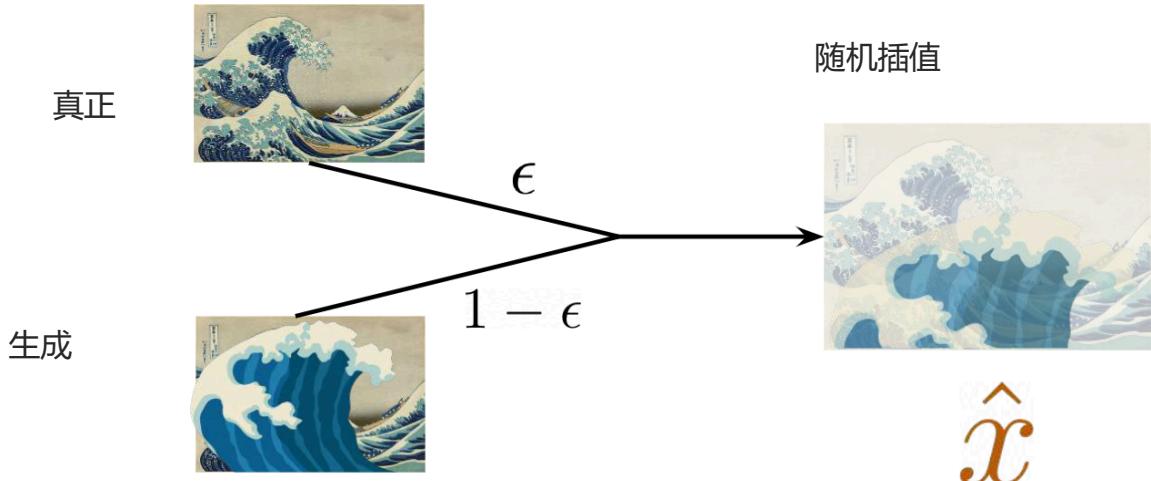
ϵ

随机插值



© deeplearning.ai

1-L 执行: 梯度惩罚



© deeplearning.ai

1-L 执行: 梯度惩罚

$$\mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2 \quad \text{正则化项}$$

© deeplearning.ai

1-L 执行: 梯度惩罚

$$\mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2 \quad \text{正则化项}$$



1-L 执行: 梯度惩罚

$$\mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2 \quad \text{正则化项}$$

↓

$$\epsilon x + (1 - \epsilon)g(z) \quad \text{插值}$$



1-L 执行: 梯度惩罚

$$\mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2 \quad \text{正则化项}$$

↓

$$\epsilon \boxed{x} + (1 - \epsilon)g(z) \quad \text{插值}$$

真正

© deeplearning.ai

1-L 执行: 梯度惩罚

$$\mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2 \quad \text{正则化项}$$

↓

$$\epsilon \boxed{x} + (1 - \epsilon) \boxed{g(z)} \quad \text{插值}$$

真正 生成

© deeplearning.ai

把它们放在一起

$$\min_g \max_c \mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) + \lambda \mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2$$



把它们放在一起

$$\min_g \max_c \boxed{\mathbb{E}(c(x)) - \mathbb{E}(c(g(z)))} + \lambda \mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2$$

使 GAN 不易出现模式折叠和渐变消失



把它们放在一起

$$\min_g \max_c [\mathbb{E}(c(x)) - \mathbb{E}(c(g(z))) + \lambda \mathbb{E}(\|\nabla c(\hat{x})\|_2 - 1)^2]$$

使 GAN 不易出现模式折叠和渐变消失

尝试使 critic 为 1-L Continuous, 使损失函数为连续且可微分



总结

- 磅码削缩和梯度惩罚是强制 1 L 连续性的方法
- 梯度惩罚往往效果更好

